

Supervised Learning

The Badges Game

+ Naoki Abe

- Eric Baum

Background:

- Pre-registered attendees at the 1994 Machine Learning Conference received a **name badge labeled with a "+" or "-"**
- The label is based only upon the name
- There are 294 examples (210 positive and 84 negative)

What function was used to generate the +/- labeling?

Training Data

- + Naoki Abe
- Myriam Abramson
- + David W. Aha
- + Kamal M. Ali
- Eric Allender
- + Dana Angluin
- Chidanand Apte
- + Minoru Asada
- + Lars Asker
- + Javed Aslam
- + Jose L. Balcazar
- Cristina Baroglio

- + Peter Bartlett
- Eric Baum
- + Welton Becket
- Shai Ben-David
- + George Berg
- + Neil Berkman
- + Malini Bhandaru
- + Bir Bhanu
- + Reinhard Blasig
- Avrim Blum
- Anselm Blumer
- + Justin Boyan

- + Carla E. Brodley
- + Nader Bshouty
- Wray Buntine
- Andrey Burago
- + Tom Bylander
- + Bill Byrne
- Claire Cardie
- + John Case
- + Jason Catlett
- Philip Chan
- Zhixiang Chen
- Chris Darken

+ Naoki Abe	- Myriam Abramson	+ David W. Aha	+ Martinch Krikis	+ Martin Kummer	- Eyal Kushilevitz
+ Kamal M. Ali	- Eric Allender	+ Dana Angluin	- Stephen Kwek	+ Wai Lam	+ Ken Lang
- Chidanand Apte	+ Minoru Asada	+ Lars Asker	- Steffen Lange	+ Pat Langley	+ Mary Soon Lee
+ Javed Aslam	+ Saralabos Athanassiou	+ Jose L. Balcazar	+ Wee Sun Lee	+ Moshe Leshno	+ Long-Ji Lin
+ Timothy P. Barber	+ Michael W. Barley	- Cristina Baroglio	- Charles X. Ling	+ Michael Littman	+ David Loewenstein
+ Peter Bartlett	- Eric Baum	+ Welton Becket	- Phil Long	+ Wolfgang Maass	- Bruce A. MacDonald
- Shai Ben-David	+ George Berg	+ Neil Berkman	+ Rich MacIain	- Sridhar Mahadevan	- J. Jeffrey Mahoney
+ Malini Bhandaru	+ Bir Bhanu	+ Reinhard Blasig	+ Yishay Mansour	+ Mario Marchand	- Shaul Markovitch
- Avrim Blum	- Anselm Blumer	+ Justin Boyan	- Oded Maron	+ Maja Mataric	+ David Mathias
+ Carla E. Brodley	+ Nader Bshouty	- Wray Buntine	+ Toshiyasu Matsushina	- Stan Matwin	- Eddy Mayoraz
- Andrey Burago	+ Tom Bylander	+ Bill Byrne	- R. Andrew McCallum	- L. Thorne McCarty	- Alexander M. Meystel
- Claire Cardie	+ Richard A. Caruana	+ John Case	+ Michael A. Meystel	- Steven Minton	+ Nina Mishra
+ Jason Catlett	+ Nicolo Cesa-Bianchi	- Philip Chan	+ Tom M. Mitchell	+ Dunja Mladenic	+ David Montgomery
+ Mark Changizi	+ Pang-Chieh Chen	- Shixiang Chen	- Andrew W. Moore	+ Johanne Morin	+ Hiroshi Motoda
+ Wan P. Chiang	- Steve A. Chien	+ Jeffery Clouse	- Stephen Muggleton	+ Patrick M. Murphy	- Sreerama K. Murthy
+ William Cohen	+ David Cohn	- Clare Bates Congdon	+ Filippo Neri	- Craig Nevill-Manning	- Andrew Y. Ng
- Antoine Cornuejols	+ Mark W. Craven	+ Robert P. Daley	+ Nikolay Nikolaev	- Steven W. Norton	+ Joseph O'Sullivan
+ Lindley Darden	- Chris Darden	+ Shaakar Dasgupta	+ Dan Oblinger	+ Jong-Moon Oh	- Arlindo Oliveira
- Brian D. Davidson	+ Michael de la Maza	+ Olivier De Vel	+ David W. Opitz	+ Sandra Panizza	+ Barak A. Pearlmutter
- Scott E. Decatur	+ Gerald F. DeJong	+ Kan Deng	- Ed Pednault	+ Jing Peng	+ Fernando Pereira
- Thomas G. Dietterich	+ Michael J. Donahue	+ George A. Drastal	+ Aurora Perez	+ Bernhard Pfahringer	+ David Pierce
+ Harris Drucker	- Chris Drummond	+ Hal Duncan	- Krishnan Pillaipakkammatt	+ Roberto Piola	+ Leonard Pitt
- Thomas Elman	+ Tapio Elomaa	+ Susan L. Epstein	+ Lorian Y. Pratt	- Armand Frieditis	+ Foster J. Provost
+ Bob Evans	- Claudio Facchinetti	+ Tom Fawcett	- J. R. Quinlan	+ John Rachlin	+ Vijay Raghavan
- Usama Fayyad	+ Aaron Feigelson	+ Nicolas Fiechter	- R. Bharat Rao	+ Priscilla Rasmussen	+ Joel Ratsaby
+ David Finton	+ John Fischer	+ Paul Fischer	+ Michael Redmond	+ Patricia J. Riddle	+ Lance Riley
+ Seth Flanders	+ Lance Fortnow	- Ameer Foued	+ Ronald L. Rivest	+ Huw Roberts	+ Dana Ron
+ Judy A. Franklin	+ Yoav Freund	+ Johannes Furnkranz	+ Robert S. Roos	+ Justinian Rosca	+ John R. Rose
+ Merrick L. Furst	+ Jean Gabriel Ganascia	+ William Gasarch	+ Dan Roth	+ James S. Royer	+ Ronitt Rubinfeld
+ Ricard Gavaldà	+ Melinda T. Gervasio	+ Yolanda Gil	- Stuart Russell	+ Lorenza Saitta	+ Yoshifumi Sakai
+ David Gillman	- Attilio Giordana	+ Kate Goetz	+ William Sakas	+ Marcos Salganicoff	- Steven Salzberg
+ Paul W. Goldberg	+ Sally Goldman	+ Diana Gordon	- Claude Sammut	+ Cullen Schaffer	+ Robert Schapire
+ Geoffrey Gordon	+ Jonathan Gratch	+ Leslie Grate	+ Mark Schwabacher	+ Michele Sebag	+ Gary M. Selzer
+ William A. Greene	+ Russell Greiner	+ Marko Grobelnik	+ Sebastian Seung	- Arun Sharma	+ Jude Shavlik
+ Tal Grossman	+ Margo Guertin	+ Tom Hancock	+ Daniel L. Silver	- Glenn Silverstein	+ Yoram Singer
+ Earl S. Harris Jr.	+ David Haussler	+ Matthias Heger	+ Mona Singh	+ Satinder Pal Singh	+ Kimmen Sjolander
+ Lisa Hellerstein	+ David Helmbold	+ Daniel Hennessy	+ David B. Skalak	+ Sean Slattery	+ Robert Sloan
+ Haym Hirsh	+ Jonathan Hodgson	+ Robert C. Holte	+ Donna Slonim	+ Carl H. Smith	+ Sonya Snedecor
+ Jiarong Hong	- Chun-Nan Hsu	+ Kazushi Ikeda	+ Von-Mun Soo	- Thomas G. Spalthoff	+ Mark Staley
+ Masayuki Inaba	- Drago Indjic	+ Nitin Indurkha	- Frank Stephan	+ Mandayam T. Suraj	+ Richard S. Sutton
+ Jeff Jackson	+ Sanjay Jain	+ Wolfgang Janko	+ Joe Suzuki	- Prasad Tadepalli	+ Hiroshi Tanaka
- Klaus P. Jantke	+ Nathalie Japkowicz	+ George H. John	- Irina Tchoumatchenko	- Brian Tester	- Chen K. Tham
+ Randolph Jones	+ Michael I. Jordan	+ Leslie Pack Kaelbling	+ Tatsuo Uemai	- Lyle H. Ungar	+ Paul Utgoff
+ Bala Kalyanasundaram	- Thomas E. Kammerer	- Grigoris Karakoulas	+ Karsten Verbeurgt	+ Paul Vitanyi	+ Xuemei Wang
+ Michael Kearns	+ Neela Khan	- Roni Khardon	+ Manfred Warmuth	+ Gary Weiss	- Sholom Weiss
+ Dennis F. Kibler	+ Jorg-Uwe Kietz	- Efin Kinber	- Thomas Wengerek	- Bradley L. Whitehall	- Alma Whitten
- Jyrki Kivinen	- Emanuel Knill	- Craig Knoblock	+ Robert Williamson	+ Janusz Wnek	+ Kenji Yamaniishi
+ Ron Kohavi	+ Pascal Koiran	+ Moshe Koppel	+ Takefumi Yamazaki	+ Holly Yanco	+ John M. Zelle
+ Daniel Kortenkamp	+ Matevz Kovacic	- Stefan Kramer	- Thomas Zeugmann	+ Jean-Daniel Zucker	+ Darko Zupanec

Test Data

- ? Shivani Agarwal
- ? Chris Callison-Burch
- ? Eric Eaton
- ? Peter Stone
- ? Matthew Taylor

Labeled Test Data

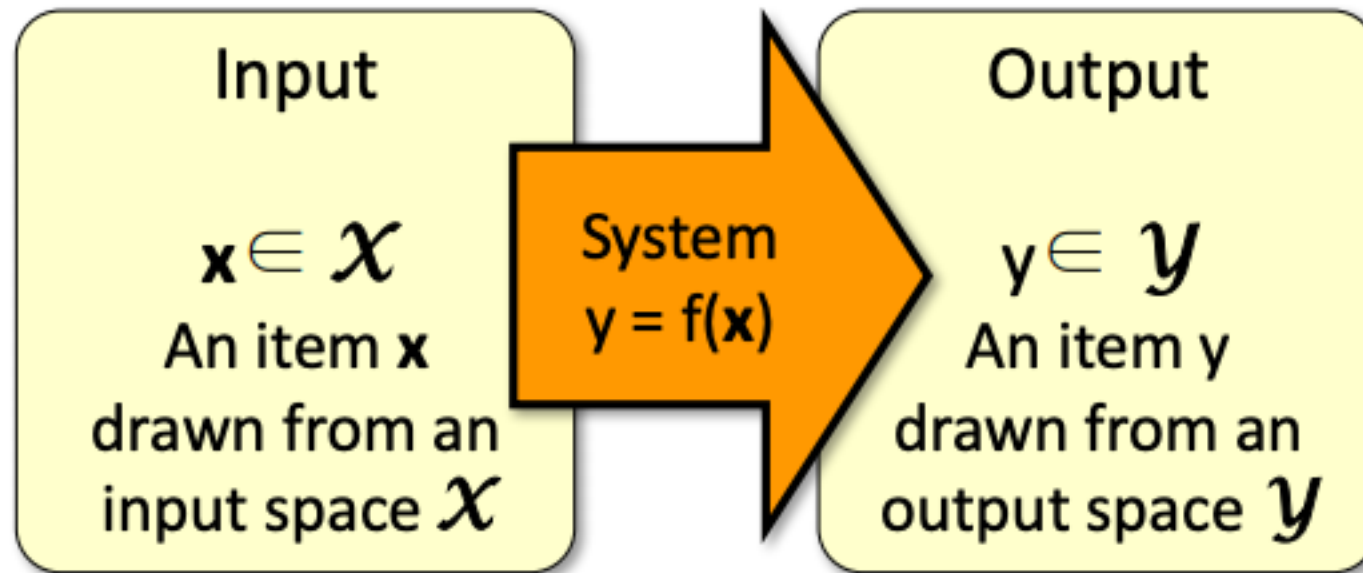
- Shivani Agarwal
- Chris Callison-Burch
- Eric Eaton
- + Peter Stone
- + Matthew Taylor

What is Learning?

- The Badges Game is an example of a key learning protocol: **supervised learning**
- First question: Are you sure you got it? Why?
- Issues:
 - Which problem was easier: prediction or modeling?
 - Representation
 - Problem setting
 - Background Knowledge
 - When did learning take place?

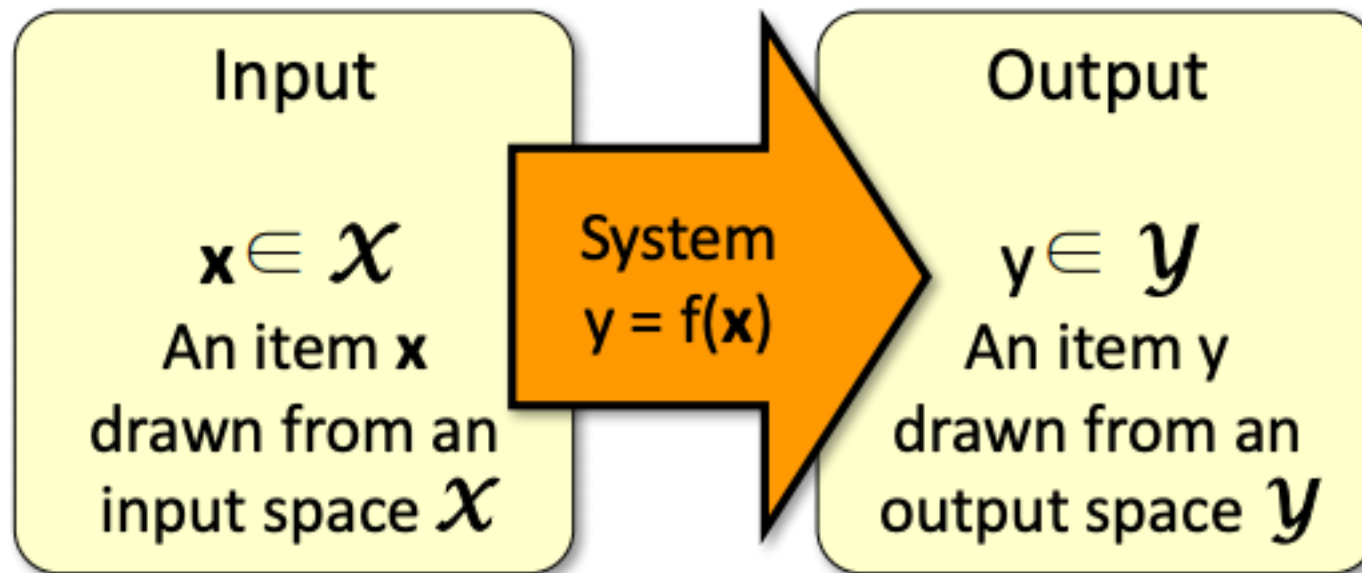
Algorithm: can you write a program that takes this data as input and predicts the label for your name?

Supervised Learning



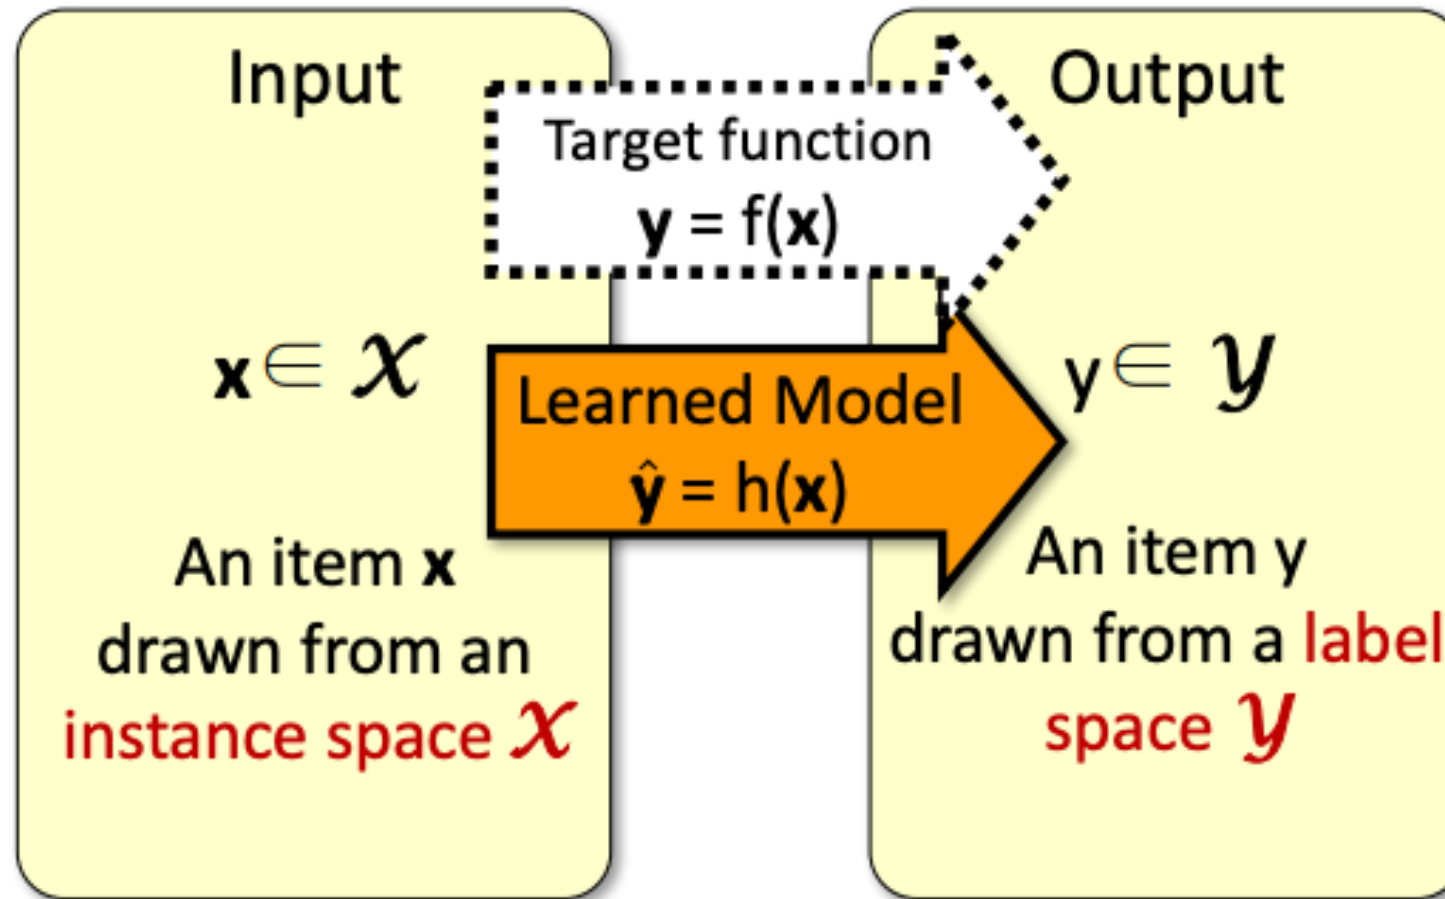
- We consider systems that apply an unknown function $f()$ to input items \mathbf{x} and return an output $y = f(\mathbf{x})$.

Supervised Learning

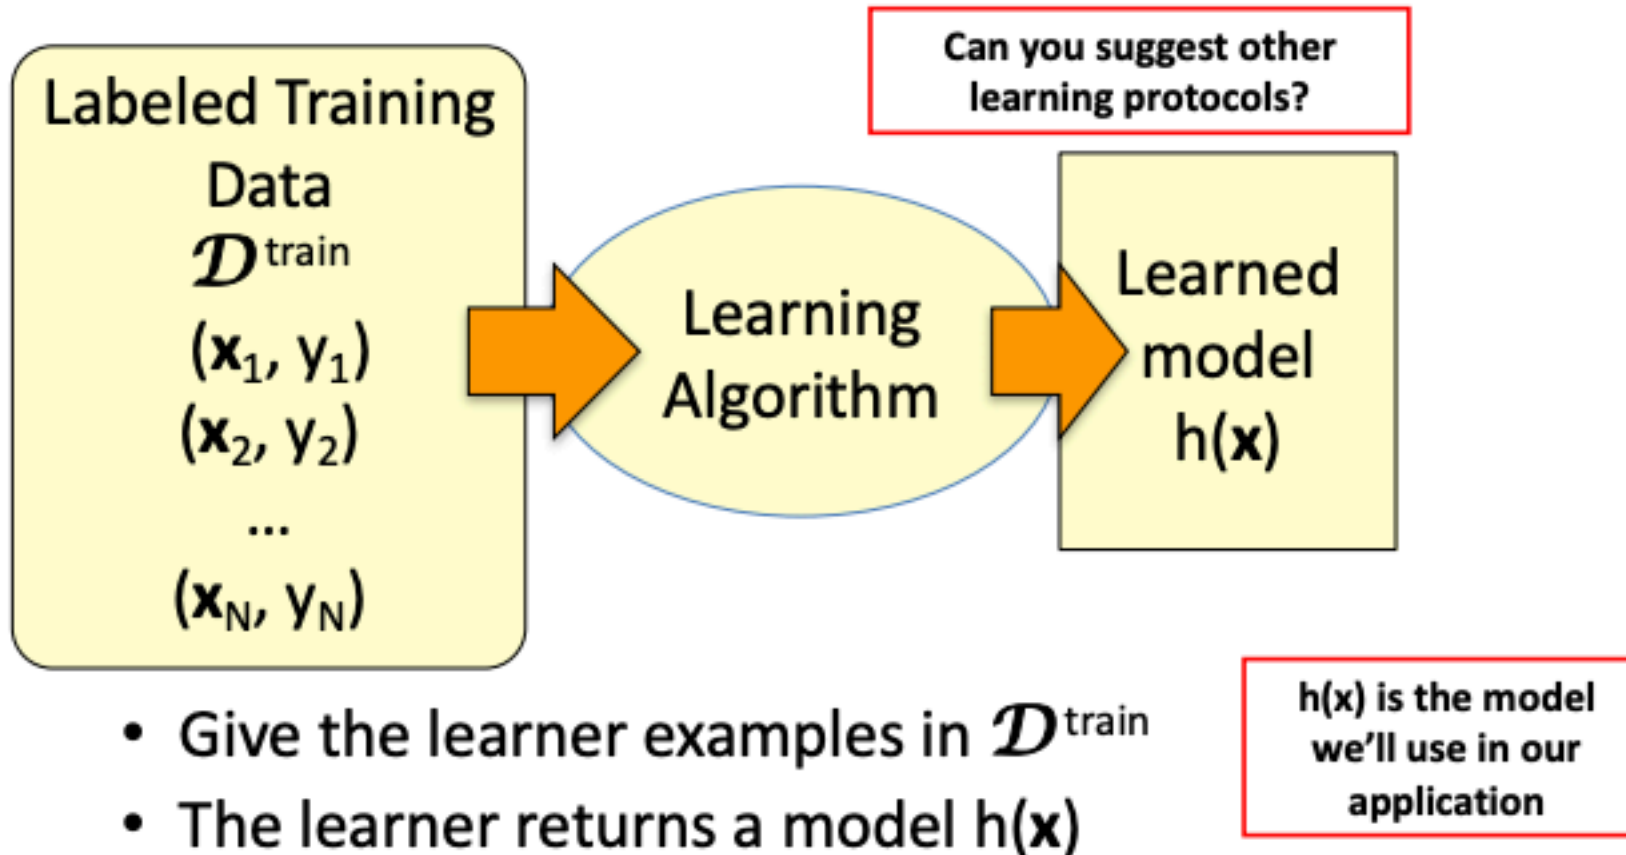


- In (supervised) machine learning, our goal is to learn a function $h()$ from examples that approximates $f()$

Supervised Learning



Supervised Learning: Training



Function Approximation

Problem Setting

- Set of possible instances \mathcal{X}
- Set of possible labels \mathcal{Y}
- Unknown target function $f : \mathcal{X} \rightarrow \mathcal{Y}$
- Set of function hypotheses $H = \{h \mid h : \mathcal{X} \rightarrow \mathcal{Y}\}$

Input: Training examples of unknown target function f

$$\{\langle \mathbf{x}_i, y_i \rangle\}_{i=1}^n = \{\langle \mathbf{x}_1, y_1 \rangle, \dots, \langle \mathbf{x}_n, y_n \rangle\}$$

Output: Hypothesis $h \in H$ that best approximates f

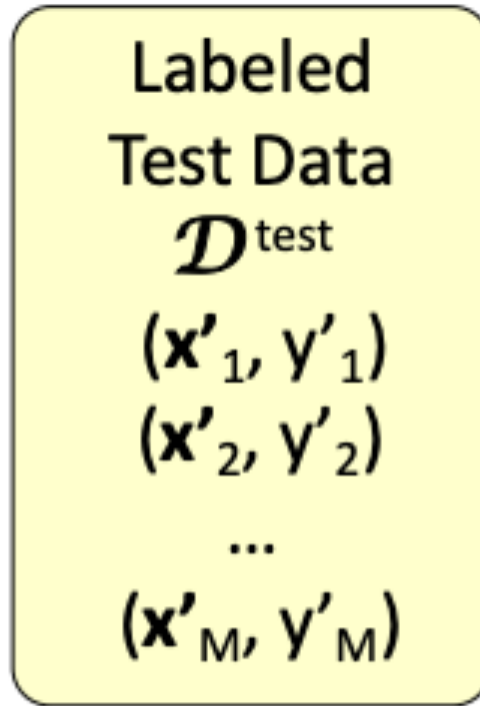
Sample Dataset

- Columns denote features X_i
- Rows denote labeled instances $\langle x_i, y_i \rangle$
- Class label denotes whether a tennis game was played

$\langle x_i, y_i \rangle$

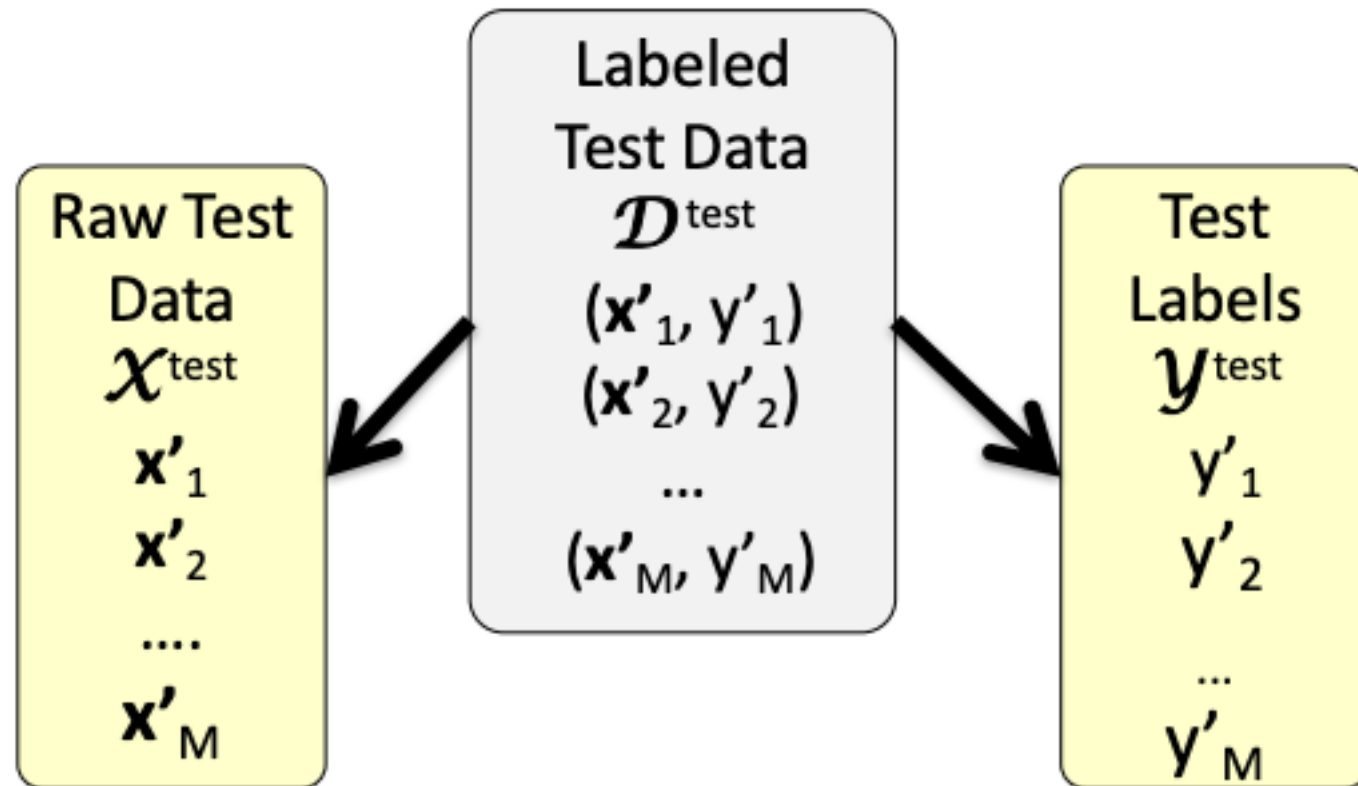
Predictors				Response
Outlook	Temperature	Humidity	Wind	Class
Sunny	Hot	High	Weak	No
Sunny	Hot	High	Strong	No
Overcast	Hot	High	Weak	Yes
Rain	Mild	High	Weak	Yes
Rain	Cool	Normal	Weak	Yes
Rain	Cool	Normal	Strong	No
Overcast	Cool	Normal	Strong	Yes
Sunny	Mild	High	Weak	No
Sunny	Cool	Normal	Weak	Yes
Rain	Mild	Normal	Weak	Yes
Sunny	Mild	Normal	Strong	Yes
Overcast	Mild	High	Strong	Yes
Overcast	Hot	Normal	Weak	Yes
Rain	Mild	High	Strong	No

Supervised Learning: Testing



- Reserve some labeled data for testing

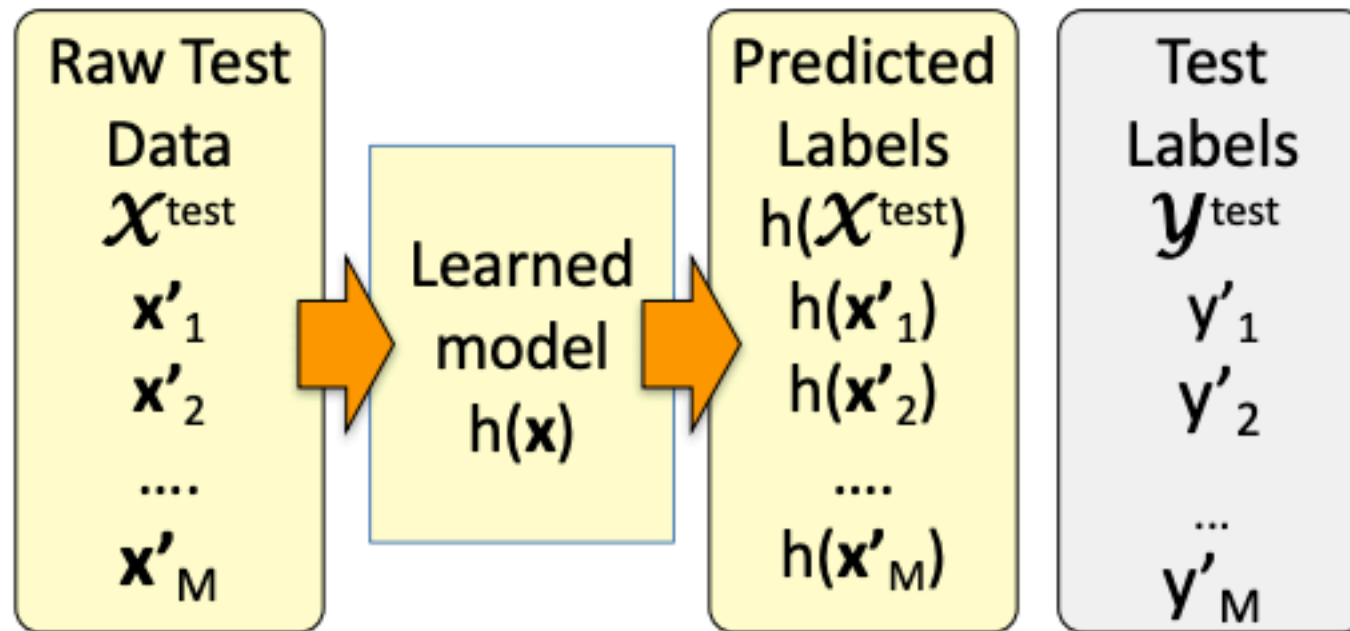
Supervised Learning: Testing



Supervised Learning: Testing

- Apply the model to the raw test data
- Evaluate by comparing predicted labels against the test labels

Can you use the test data otherwise?



Supervised Learning: Examples

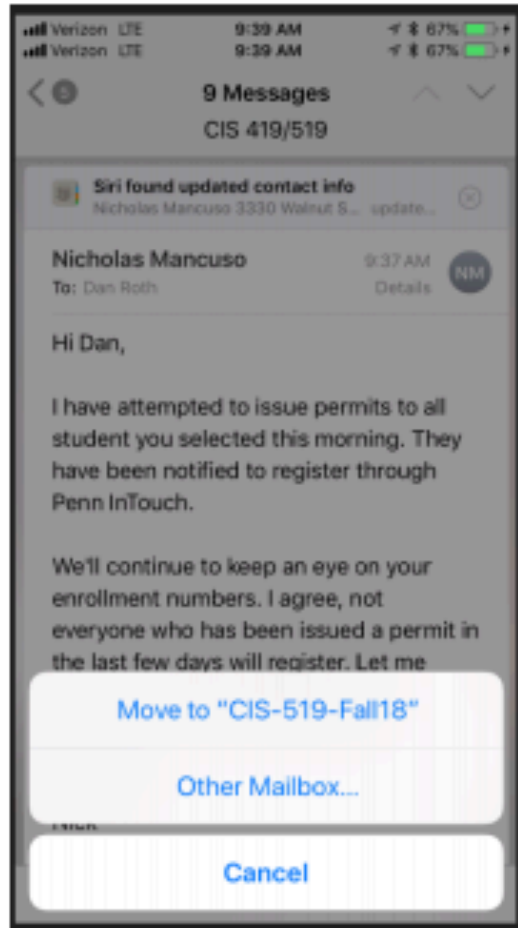
- Disease diagnosis
 - x : Properties of patient (symptoms, lab tests)
 - f : Disease (or maybe: recommended therapy)
- Part-of-Speech tagging
 - x : An English sentence (e.g., The can will rust)
 - f : The part of speech of a word in the sentence
- Face recognition
 - x : Bitmap picture of person's face
 - f : Name the person (or maybe: a property of)
- Automatic Steering
 - x : Bitmap picture of road surface in front of car
 - f : Degrees to turn the steering wheel

Many problems that do not seem like classification problems can be decomposed into classification problems.

Key Issues in Machine Learning

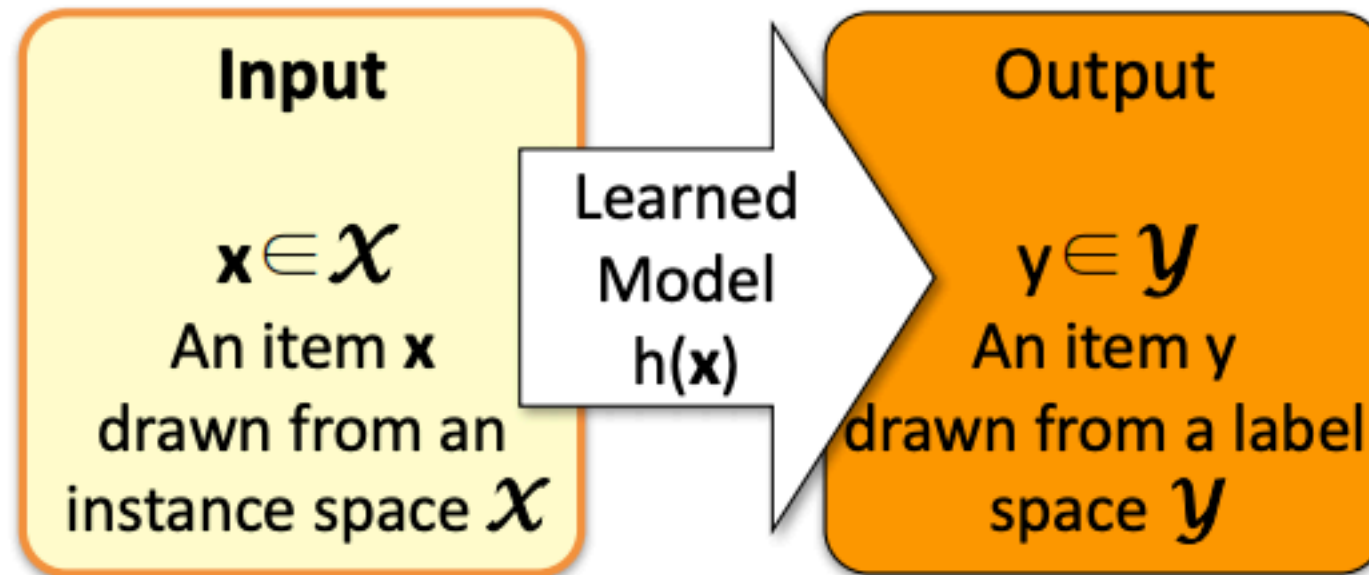
- Modeling
 - How to formulate application problems as machine learning problems?
 - How to represent the data?
 - Learning Protocols (**where is the data & labels coming from?**)
- Representation
 - What **functions** should we learn (hypothesis spaces) ?
 - How to map raw **input** to an instance space?
 - Any rigorous way to find these? Any general approach?
- Algorithms
 - What are good algorithms?
 - How do we define success?
 - Generalization vs. overfitting
 - The computational problem

Using Supervised Learning



- What is our instance space?
 - What kind of features are we using?
- What is our label space?
 - What kind of learning task are we dealing with?
- What is our hypothesis space?
 - What kind of functions (models) are we learning?
- What learning algorithm do we use?
 - How do we learn the model from the labeled data?
- What is our loss function/evaluation metric?
 - How do we measure success? What drives learning?

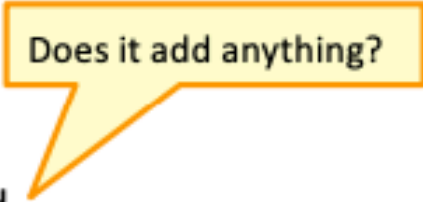
1. The Instance space \mathcal{X}



- Designing an appropriate instance space \mathcal{X} is crucial for how well we can predict \mathbf{y} .

1. The Instance space \mathcal{X}

- When we apply machine learning to a task, we first need to define the instance space \mathcal{X} .
- Instances $x \in \mathcal{X}$ are defined by features:
 - **Boolean features:**
 - Is there a folder named after the sender?
 - Does this email contains the word 'class'?
 - Does this email contains the word 'waiting'?
 - Does this email contains the word 'class' and the word 'waiting'?
 - **Numerical features:**
 - How often does 'learning' occur in this email?
 - What long is email?
 - How many emails have I seen from this sender over the last day/week/month?
 - **Bag of tokens**
 - Just list all the **tokens** in the input




Does it add anything?

What's \mathcal{X} for the Badges game?

- Possible features:
 - Gender
 - Name's country-of-origin
 - Length of their first or last name
 - Does the name contain letter 'x'?
 - How many vowels does their name contain?
 - Is the n-th letter a vowel?
 - Does the name have the same number of vowels and consonants?

\mathcal{X} as a Vector Space


- \mathcal{X} is an N-dimensional vector space (e.g. \mathbb{R}^N)
 - Each dimension = one feature.
- Each \mathbf{x} is a **feature vector** (hence the boldface \mathbf{x}).
- Think of $\mathbf{x} = [x_1 \dots x_N]$ as a point in \mathcal{X} :



Good Features are Essential

- The choice of features is **crucial** for how well a task can be learned
 - In many application areas (language, vision, etc.), a lot of work goes into designing suitable features
 - This requires domain expertise
- Think about the badges game – what if you were focusing on visual features?
- We can't teach you what specific features to use for your task
 - But we will touch on some general principles

1. The Label space \mathcal{Y}



- The label space \mathcal{Y} determines *what kind of supervised learning task* we are dealing with

Supervised Learning Task I

- Output labels $y \in Y$ are categorical:
 - **Binary classification**: Two possible labels
 - **Multi-class classification**: k possible labels
- Output labels $y \in Y$ are **structured objects** (sequences of labels, parse trees, etc.)
- Structure learning




Before
I met with him before leaving for Paris
on Thursday. Be_Included

Supervised Learning Task II

- Output labels $y \in Y$ are numerical:
 - Regression (linear/polynomial):
 - Labels are continuous-valued
 - Learn a linear/polynomial function $f(x)$
 - Ranking:
 - Labels are ordinal
 - Learn an ordering $f(x_1) > f(x_2)$ over input

3. The Model $h(x)$



- We need to choose what *kind* of model we want to learn

A Learning Problem



Example	x_1	x_2	x_3	x_4	y
1	0	0	1	0	0
2	0	1	0	0	0
3	0	0	1	1	1
4	1	0	0	1	1
5	0	1	1	0	0
6	1	1	0	0	0
7	0	1	0	1	0

Can you learn this function? What is it?

Hypothesis Space

Complete Ignorance:

There are $2^{16} = 65536$ possible functions over four input features.

We can't figure out which one is correct until we've seen every possible input-output pair.

After observing seven examples we still have 2^9 possibilities for f

Is Learning Possible?

Example	X1	X2	X3	X4	Y
1	0	0	0	0	?
2	0	0	0	1	?
	0	0	1	0	0

- There are $|Y|^{|X|}$ possible functions $f(x)$ from the instance space X to the label space Y .
- Learners typically consider **only a subset of the functions from X to Y** , called the hypothesis space H . $H \subseteq |Y|^{|X|}$

	1	1	0	1	?
	1	1	1	0	?
16	1	1	1	1	?

General Strategies for Machine Learning

- Develop flexible hypothesis spaces:
 - Decision trees, neural networks, nested collections.
 - **Constraining the hypothesis space is done algorithmically**
- Develop representation languages for restricted classes of functions:
 - Serve to limit the expressivity of the target models
 - E.g., Functional representation (n-of-m); Grammars; linear functions; stochastic models;
 - **Get flexibility by augmenting the feature space**
- In either case:
 - Develop algorithms for finding a hypothesis in our hypothesis space, that fits the data
 - And hope that they will generalize well

Key Issues in Machine Learning

- Modeling
 - How to formulate application problems as machine learning problems?
 - How to represent the data?
 - Learning Protocols (**where is the data & labels coming from?**)
- Representation
 - What **functions** should we learn (hypothesis spaces) ?
 - How to map raw **input** to an instance space?
 - Any rigorous way to find these? Any general approach?
- Algorithms
 - What are good algorithms?
 - How do we define success?
 - Generalization vs. overfitting
 - The computational problem